

Problem Set 1

Lars Lien Ankile

SECTION LEADER + TIME

Problem 1.

- a) This method of selected gives no guarantees about the diverseness of the sample and whether or not this sample will be representative for the whole. E.g. if all the smallest firms happens to come before the larger firms in the phonebook, then the data collected might be very misleading.
- b) The readers of Prevention magazine is not unlikely a self-selected group of people who care a lot about their health in the first place, and that is why they chose to subscribe to the magazine. This will in turn result in a very unrepresentative sample for the population at large.
- c) Even though the clinics are accessible to anyone, that doesn't mean that all homeless people would actually go to a clinic. I find it more likely that the homeless who are actually suffering from mental illness are more likely to go to a clinic than those who are not suffering.

Problem 2.

- a) It can be dangerous to mix up correlation with causation, i.e. just because two events often are observed together, that doesn't necessarily mean that one event is making the other happen. This headline makes it sound like people who are more optimistic suffer less from cardiovascular disease *because* they're more positive. If that is true is unknown and the only thing we can conclude from this is that people who are optimistic also *happen* to also suffer less from cardiovascular disease, for *some* reason.
- b) To have a study actually say something about causation and not only correlation, one would have to design a study that controls for all spurious causes that could impact both positivity and cardiovascular disease. Somehow, if one could take a random sample of the population, divide it into two groups with roughly similar people in both groups, and then make one group think optimistically and make the other think pessimistically, while trying to control for any other effect, one could potentially say a little bit more about the effect, I'd imagine.
- c) Just because someone is on average less likely to have some event occur it doesn't mean that all optimistic people have lower risk of cardiovascular events than all pessimistic people. Surely there will be pessimistic people who just happen to have a great cardiovascular system and will live to a hundred, while there are optimistic smokers who'll die from a hearth attack tomorrow.

Problem 3.

- a) Both graphs show the distribution of the ages that women had their first child for a given year. The year is 1980 for the first and 2016 for the second. In 1980, for example, we see that some, but very few people had their first child at 14, while having one's first child was the most common at roughly 19. The graph is not symmetric as it has a long tail to the right of the distribution and is therefore right-skewed. The graph for 2016 shows a shift of some of the population to the right. Also, there's two distinct tops on the graph. This graph also has a tail to the right which makes it right-skewed too, but not as much as the former one, though.
- b)
- i. I think that the change in the distribution can be best explained by an increase in the percentage of women who take higher education.
 - ii. There's a top at 19 as in the graph for 1980, but now there's also one at age 28. This suggests to me that many people self-select into two groups, presumably based on whether they take higher education or not. The second peak on the graph comes from people who have taken education and waited until they are more established before they have their first child.

Problem 4.

a)

```
#load the data
load("datasets/gun_deaths.Rdata")

#numerical summaries
# Females
fem_mean = mean(gun_deaths$age[gun_deaths$sex == 'F'], na.rm = TRUE)
fem_median = median(gun_deaths$age[gun_deaths$sex == 'F'], na.rm = TRUE)

sprintf('For females the mean age is %.2f and the median age is %.2f', fem_mean, fem_median)

## [1] "For females the mean age is 43.70 and the median age is 44.00"

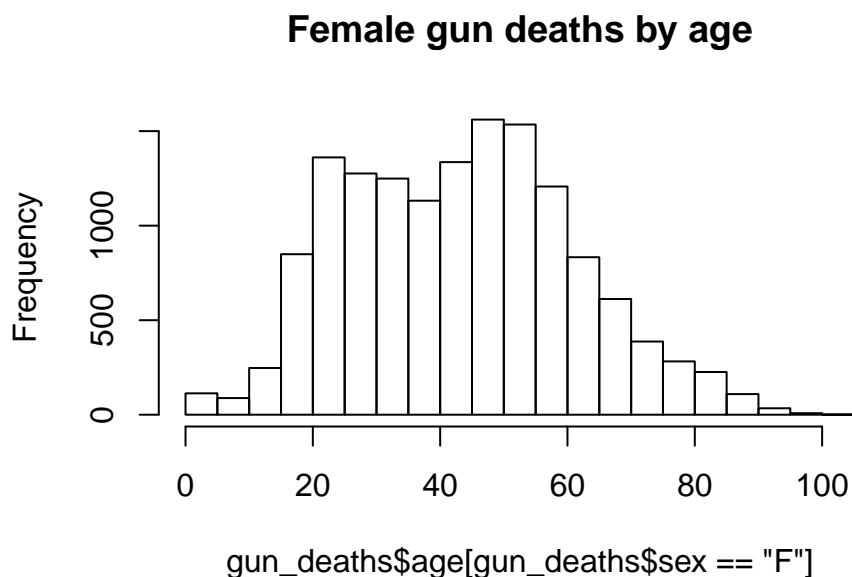
# Males
male_mean = mean(gun_deaths$age[gun_deaths$sex == 'M'], na.rm = TRUE)
male_median = median(gun_deaths$age[gun_deaths$sex == 'M'], na.rm = TRUE)

sprintf('For males the mean age is %.2f and the median age is %.2f', male_mean, male_median)

## [1] "For males the mean age is 43.88 and the median age is 41.00"


#graphical summaries

# Females
hist(gun_deaths$age[gun_deaths$sex == 'F'], main = "Female gun deaths by age")
```



```
# Males
hist(gun_deaths$age[gun_deaths$sex == 'M'], main = "Male gun deaths by age")
```

Male gun deaths by age



`gun_deaths$age[gun_deaths$sex == "M"]`

From the above we see that the female distribution is slightly left-skewed, while the male distribution is right-skewed (mean is smaller than median in the first case and mean is greater than median in the second case). The difference between the mean and median is relatively small in the case for the females, though. From the histograms we see that the most common age to be killed at as a man is right after 20, which is what one might expect. A little more surprising, to me at least, is that the highest rate of gun deaths for females happen between 40 and 60.

b)

```
table(gun_deaths$intent)
```

```
##
##   Accidental   Homicide   Suicide Undetermined
##         1639        35176        63175         807
```

We see that suicide is by far the most common category for cause of gun death. Homicides also contribute a lot, while accidental deaths and undetermined deaths are much less common.

c)

```
# Get info
table(gun_deaths$intent, gun_deaths$police)
```

```
##
##           No    Yes
## Accidental  1639    0
## Homicide   33774  1402
## Suicide    63175    0
## Undetermined 807    0
```

```
# Calculate proportion
total = 1402 + 33774
1402 / (total)
```

```
## [1] 0.03985672
```

```
total
```

```
## [1] 35176
```

1402 out of a total of 35176, or about 4.5%, of the gun deaths classified as homicides was caused by police intervention.

d)

```
table(gun_deaths$month)
```

```
##
```

```
##      1      2      3      4      5      6      7      8      9     10     11     12
```

```
## 8273 7093 8289 8455 8669 8677 8989 8783 8508 8406 8243 8413
```

```
num_spring = length(gun_deaths$month[gun_deaths$month >= 3 & gun_deaths$month < 6])
```

```
num_summer = length(gun_deaths$month[gun_deaths$month >= 6 & gun_deaths$month < 9])
```

```
num_fall = length(gun_deaths$month[gun_deaths$month >= 9 & gun_deaths$month < 12])
```

```
num_winter = length(gun_deaths$month[gun_deaths$month >= 12 | gun_deaths$month < 3])
```

```
sprintf("Gun death per season: Spring: %d, Summer: %d, Fall: %d, Winter: %d.", num_spring, num_
```

```
## [1] "Gun death per season: Spring: 25413, Summer: 26449, Fall: 25157, Winter: 23779."
```

The summer was the season with the most gun deaths with its 26449 gun deaths.

e)

```
deaths.2012 = gun_deaths[gun_deaths$year == 2012 & gun_deaths$education %in% c("HS/GED", "Some  
table(deaths.2012$race)
```

```
##
```

```
##      Asian/Pacific Islander      Black
```

```
##              372              5207
```

```
##      Hispanic Native American/Native Alaskan
```

```
##              1651              197
```

```
##      White
```

```
##      18121
```

```
# Missing values tho?
```

```
white_deaths = nrow(deaths.2012[deaths.2012$race == 'White', ])
```

```
total = nrow(deaths.2012)
```

```
sprintf("Out of the total of %d deaths in 2012 among people with at least high school education
```

```
## [1] "Out of the total of 25548 deaths in 2012 among people with at least high school educat
```

```
sprintf("I.e. %.2f%% of the people killed were white.", white_deaths * 100 / total)
```

```
## [1] "I.e. 70.93% of the people killed were white."
```

As we can see from the above, the majority of the people killed were white with roughly 71% being white.

f)

```
table(gun_deaths$race, gun_deaths$intent)
```

```
##
##               Accidental Homicide Suicide Undetermined
## Asian/Pacific Islander          12      559      745          10
## Black                          328     19510     3332         126
## Hispanic                       145      5634     3171          72
## Native American/Native Alaskan    22       326      555          14
## White                         1132      9147     55372         585
```

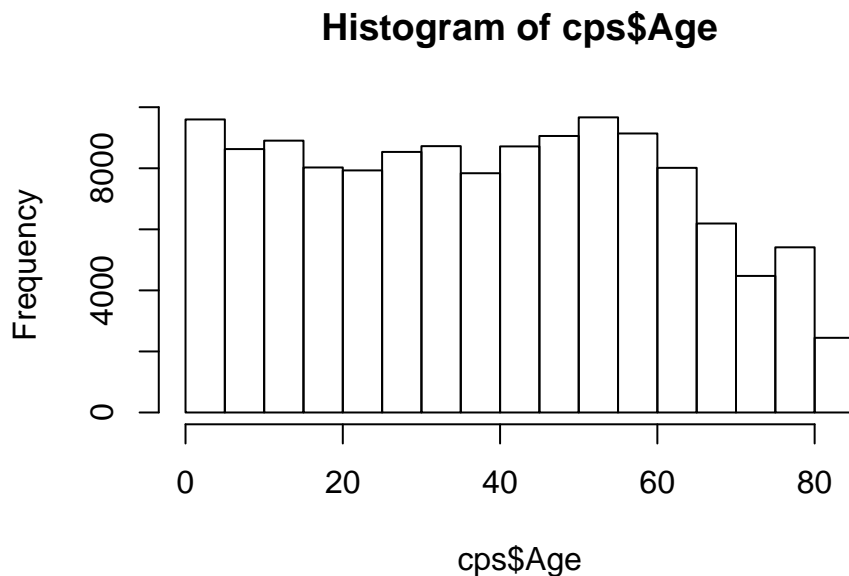
From the above I notice especially two glaring things from the data. First, for black people, there are almost 6 times as many homicides as there are suicides. Second, for white people, there are almost 6 times as many suicides than homicides, i.e. the exact opposite of black people.

Problem 5.

a) i.

```
#load the data
cps = read.csv("datasets/CPSData.csv", header = TRUE)

#explore age
hist(cps$Age)
```



```
#explore sex
table(cps$Sex)
```

```
##
## Female    Male
## 67481    63821
```

```
#explore race
table(cps$Race)
```

```
##
## American Indian      Asian      Black      Multiracial
##           1433          6520      13913           2897
## Pacific Islander      White
##           618          105921
```

From the age histogram above we see that there's a relatively even distribution of people in the ages between 0 and 60. After 60 there's a dropoff of number of people, which makes sense. From the first table above we see that there's a little more women than men in the sample, which makes sense given that there's more women in society at large as well, presumably because women live longer. When it comes to race, we see that there's by far most white people with almost 106 thousand people. Black people comes second with almost 14 thousand, which is significantly less.

ii. One thing I find slightly unusual about the age of the people is that there's a relatively large

spike in the number of 75-80 year-olds as compared to those 70-75 and 80-85. One would suspect a more “linear” descent. However, these people are the people who would’ve been born after the second world war, which makes sense because we know that there was a big baby boom after the war.

b)

```
#numerical summaries
```

```
summary(cps$PeopleInHousehold)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      1.000   2.000   3.000   3.284   4.000   15.000
```

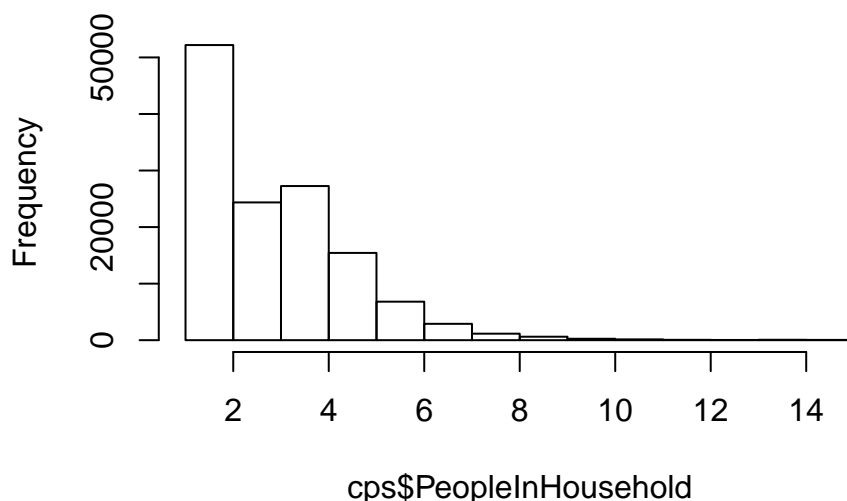
```
IQR(cps$PeopleInHousehold)
```

```
## [1] 2
```


```
#graphical summaries
```

```
hist(cps$PeopleInHousehold)
```

Histogram of cps\$PeopleInHousehold



```
boxplot(cps$PeopleInHousehold, horizontal = TRUE)
```



The median number of people in a household is 3, while the mean is 3.284, which means our data is right-skewed. This makes sense to me because most people are either living alone or in a family of 2, 3 or 4 people. Still, there are people

who have much larger households. This is confirmed by the histogram. The boxplot confirms that there are a lot of outliers who have a lot of people in their households.

c)

```
summary(cps$Citizenship)
```

```
##      Citizen, Native Citizen, Naturalized      Non-Citizen
##      116639                7073                7590
```

```
nrow(cps[cps$Citizenship == "Citizen, Native", ]) / nrow(cps)
```

```
## [1] 0.8883261
```

The vast majority of people are native citizens, with 89% of people belonging to that category. The remaining 11% is more or less evenly divided between naturalized citizens and non-citizens.

d)

```
prop.table(table(cps$Race, cps$Hispanic), 1)
```

```
##
##              0              1
## American Indian 0.78785764 0.21214236
## Asian          0.98266871 0.01733129
## Black          0.95536549 0.04463451
## Multiracial     0.84535727 0.15464273
## Pacific Islander 0.87540453 0.12459547
## White          0.84204265 0.15795735
```

From the above we see that American Indians, Multiracials, and Whites are the races where at least 15% identify as hispanic.

e)

```
table(cps$Age, cps$Married)
```

```
##
##      Divorced Married Never Married Separated Widowed
## 0           0         0             0           0         0
## 1           0         0             0           0         0
## 2           0         0             0           0         0
## 3           0         0             0           0         0
## 4           0         0             0           0         0
## 5           0         0             0           0         0
## 6           0         0             0           0         0
## 7           0         0             0           0         0
## 8           0         0             0           0         0
## 9           0         0             0           0         0
## 10          0         0             0           0         0
## 11          0         0             0           0         0
## 12          0         0             0           0         0
## 13          0         0             0           0         0
```

##	14	0	0	0	0	0
##	15	6	19	1753	16	1
##	16	10	8	1721	8	4
##	17	5	19	1729	10	1
##	18	6	21	1559	9	1
##	19	8	47	1448	14	0
##	20	6	56	1328	5	3
##	21	11	118	1378	16	2
##	22	22	187	1306	17	4
##	23	18	266	1333	18	3
##	24	27	323	1245	30	2
##	25	31	411	1132	29	1
##	26	54	484	1064	32	9
##	27	60	587	975	27	8
##	28	80	761	857	35	3
##	29	75	814	731	21	4
##	30	82	930	791	46	5
##	31	96	922	704	37	3
##	32	129	1037	573	41	10
##	33	135	1089	534	42	4
##	34	130	1000	483	34	6
##	35	161	1047	433	62	13
##	36	142	1022	443	48	8
##	37	138	1016	325	43	9
##	38	191	958	323	46	12
##	39	161	1060	264	45	12
##	40	204	1034	276	46	11
##	41	211	1126	266	54	16
##	42	242	1149	257	50	13
##	43	256	1218	270	50	25
##	44	232	1199	255	49	29
##	45	272	1115	278	58	26
##	46	251	1106	243	42	23
##	47	241	1090	240	56	20
##	48	290	1198	228	45	30
##	49	312	1339	242	60	36
##	50	341	1290	243	47	45
##	51	312	1269	256	53	41
##	52	326	1285	231	41	52
##	53	350	1275	266	53	50
##	54	357	1228	216	55	56
##	55	329	1279	187	36	64
##	56	309	1312	191	38	85
##	57	324	1211	186	41	65
##	58	312	1232	200	41	89
##	59	286	1194	161	37	80
##	60	279	1195	158	34	80
##	61	297	1142	165	34	97

##	62	252	1097	136	28	82
##	63	260	1065	118	28	125
##	64	268	1011	119	22	99
##	65	251	1057	103	20	138
##	66	228	1073	101	23	152
##	67	173	837	73	14	130
##	68	179	741	71	11	128
##	69	140	705	52	13	152
##	70	165	779	59	12	180
##	71	135	655	56	18	167
##	72	107	624	38	12	160
##	73	98	583	27	8	180
##	74	98	533	39	9	163
##	75	82	456	28	10	187
##	76	63	431	28	8	199
##	77	69	424	26	7	172
##	78	73	372	22	4	188
##	79	54	352	36	3	216
##	80	219	1269	91	15	1070
##	85	120	757	102	11	1456

From the above table we see that the amount of married people increases with age up to a certain point, i.e. around 50. After this point we see a lot of divorces in the years following, but the number of divorces is decreasing as people get older, which makes sense, I think. Also, the number of widowed people increase as people age, which also makes a lot of sense. What I found surprising, though, was the amount of 15-year olds that were divorced, a whopping 6 people, not to mention the one 15-year old who is widowed and the 16 who are separated.

Problem 6.

```
#load quantmod package
library(quantmod)

#load AAPL and MSFT data
getSymbols("AAPL", from = "2018-01-01", to = "2019-07-22")

## [1] "AAPL"

getSymbols("MSFT", from = "2018-01-01", to = "2019-07-22")

## [1] "MSFT"


#obtain adjusted closing prices
aapl.closing = Ad(AAPL)
msft.closing = Ad(MSFT)

#obtain daily volume
aapl.volume = Vo(AAPL)
msft.volume = Vo(MSFT)


#obtain daily returns
aapl.return = as.numeric(dailyReturn(Ad(AAPL)))
msft.return = as.numeric(dailyReturn(Ad(MSFT)))
```

a)

```
#plot prices
plot(aapl.closing, col = "blue", ylim = c(50, 300))
```



```
lines(msft.closing, col = "red")
```



Both companies had a positive price change over the period. It looks like Microsoft had a slightly better return. One big difference between them is that Apple looks much more volatile and more risky. It grew more from Jan '18 to Oct '18 than Microsoft, but Microsoft didn't have the drastic drop in price between Oct '18 and Jan '19 as Apple had. There's a similar effect in May '19. Microsoft might've been a better investment in this period.

b)

```
#example: summary of aapl.closing
summary(as.numeric(aapl.closing))
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  140.1   169.2   183.1   184.2   198.7   227.8
```

```
#return for aapl
```

```
100 * (as.numeric(aapl.closing)[length(aapl.closing)] - as.numeric(aapl.closing)[1]) / as.numeric(aapl.closing)[1]
```

```
## [1] 20.34641
```

```
#return for msft
```

```
100 * (as.numeric(msft.closing)[length(msft.closing)] - as.numeric(msft.closing)[1]) / as.numeric(msft.closing)[1]
```

```
## [1] 62.99628
```

From this we see that Microsoft had a much better return than Apple in this period. The fact that the scale on the y-axis is linear and Microsoft's price started out lower might be a little misleading in the graph above.

c)

```
#sd for aapl
sd(aapl.closing)
```

```
## [1] 19.54567
```

```
#sd for msft  
sd(msft.closing)
```

```
## [1] 13.54294
```

From the above we see that Microsoft is indeed the least volatile of the two stocks.

d)

```
# Min and max for Apple  
sprintf('Apple: Max value: %.2f on %s, and min value: %.2f on %s',  
        max(aapl.closing),  
        index(aapl.closing[which.max(aapl.closing)]),  
        min(aapl.closing),  
        index(aapl.closing[which.min(aapl.closing)]))
```

```
## [1] "Apple: Max value: 227.84 on 2018-10-03, and min value: 140.09 on 2019-01-03"
```

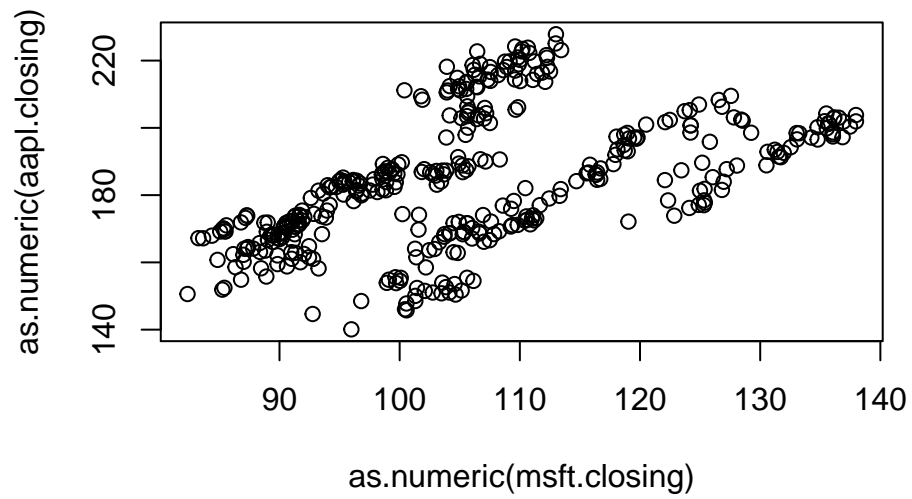
```
# Min and max for Microsoft  
sprintf('Microsoft: Max value: %.2f on %s, and min value: %.2f on %s',  
        max(msft.closing),  
        index(msft.closing[which.max(msft.closing)]),  
        min(msft.closing),  
        index(msft.closing[which.min(msft.closing)]))
```

```
## [1] "Microsoft: Max value: 137.97 on 2019-07-12, and min value: 82.35 on 2018-02-08"
```

See above.

- e)
 - i. From the graph in a) above, it seems like Microsoft and Apple are pretty correlated, at least when it comes to the direction of the price movement, i.e. when apple goes up, microsoft does too, and vice versa. However, we see that Apple's movements are much more extreme than Microsoft's.
 - ii. Used cor to calculate how correlated the stock prices are. See below.
 - iii. The correlation between the stock prices is at 0.5153. This is a positive number which means that there is a positive correlation between them. This means that the price of one tends to move in the same direction as the other over time. However, we see that the correlation is at roughly 0.5, which means that this is by no means a perfect correlation, i.e. there is a lot of variability between the stocks, too.

```
#graphical summary (part i.)  
plot(as.numeric(msft.closing), as.numeric(aapl.closing))
```



```
#numerical summary (part ii.)  
cor(as.numeric(msft.closing), as.numeric(apl.closing))
```

```
## [1] 0.5152774
```